

AL JAZEERA MEDIA INSTITUTE · PROFESSIONAL TRAINING

Post-Production with AI

DAY 1 TAKEAWAYS — Introduction to AI Video & Audio Generation

Inside this guide: the 2026 model cheat sheet · Runway field tips · the Google stack workflow (Gemini → Nano Banana → Veo) · a ready-to-use prompt library · audio generation playbook · glossary & do/don't checklist.

Trainer **Abdullah Abudiak**

Founder & CEO, The Aigency — Dubai

Instagram @a.abudiak · [linkedin.com/in/abdullahabudiak](https://www.linkedin.com/in/abdullahabudiak)

1 • The 2026 Model Cheat Sheet

Updated June 2026. The landscape moves monthly — judge models by what they solve for you, not by launch hype.

Model	Maker	Why it matters	Reach for it when...
Runway Gen-4.5	Runway	Text-to-video with native audio, multi-shot sequencing, character-consistent clips up to ~1 minute. Image-to-video (first-frame) added Jan 2026.	You need a professional pipeline: generate, direct, and edit in one place.
Runway Gen-4 / Turbo	Runway	Image-to-video workhorse with References for stable characters, objects and environments across shots. Turbo is the cheap, fast iteration variant.	Drafting and iterating — burn Turbo credits on tests, save Gen-4.5 for finals.
Veo 3.1	Google	Cinematic generation with native audio and dialogue, vertical 9:16, 1080p/4K upscale, and Ingredients-to-Video for combining reference images into one scene.	You're inside the Google ecosystem (Gemini app, Flow, Vids) or need vertical social video.
Nano Banana 2 & Pro	Google	Gemini-3 image models. Pro: 1K/2K/4K, crisp typography, controlled composition. NB2: web-grounded, fast, strong conversational multi-turn editing.	Keyframes, ingredient images, posters, anything with on-image text.
Sora 2 Pro	OpenAI	Long-form, narrative-driven text-to-video; launched aggressively in early 2026 and strong at story beats from a single prompt.	Concept films and narrative previz where one long prompt should carry a story.
Kling 3.0	Kuaishou	Native 4K @ 60fps — highest native res/frame-rate on the market. Multi-shot storyboards: up to 6 shots in a 15-second clip, with reference-based character consistency.	Maximum image quality, sports-style motion, or precise shot-by-shot storyboarding.

Insider note: one Runway subscription now bundles third-party models — Veo 3.1, Kling 3.0 Pro, Sora 2 Pro and more run inside the same dashboard, so you can compare models on the same prompt without juggling accounts.

2 • Runway Field Tips

What separates a team that gets one usable shot in twenty from a team that gets one in three.

Workflow

- **Start from an image, not text** — Image-to-video gives you composition control for free. Generate or shoot the first frame, then animate it — your hit rate will roughly double.
- **Lock identity with References** — Upload your character/object as a Gen-4 Reference and reuse it across every shot. This is how you keep the same presenter, product, or location through an entire sequence.
- **Iterate cheap, finish expensive** — Test ideas on Gen-4 Turbo (~50 credits per 10s), then re-run only the winning prompt on Gen-4.5 for the final.
- **Save your seeds** — When a generation works, save the seed value. Re-using it keeps the look while you adjust one variable — the fastest path to a matched series of shots.
- **Use Explore mode for batches** — Slower queue, but unlimited generations — ideal for overnight option-generation before an edit day.

Directing the shot

- **Motion brush > hoping** — Paint exactly what should move (and how much). One brushstroke replaces five prompt rewrites.
- **Speak film language** — “Slow dolly-in”, “handheld tracking shot”, “rack focus” — the model was trained on cinema and follows the grammar.
- **One scene per generation** — Don't cram a three-beat story into one prompt. Generate beats separately and cut them together — you're editors; act like it.

Fixing instead of regenerating

- **Aleph is your in-video eraser** — Remove objects, change weather, restyle a shot inside existing footage — try a fix before you burn credits on a re-roll.
- **Act-Two for performance** — Record yourself on a phone; Act-Two transfers your face, body and hand performance onto your generated character. Direction beats description.
- **Audio tab before export** — TTS, SFX and speech-to-speech live in the same workspace — rough-mix your sound inside Runway before you ever open an NLE.

3 • The Google Stack Workflow

Gemini → Nano Banana → Veo 3.1: the three-step chain we used in the hands-on session, written down so you can repeat it tomorrow.

Step 1 — Gemini as creative director

Don't write video prompts cold. Brief Gemini like a producer: give it the story, audience and duration, and ask for a shot list with one prompt per shot. Then ask it to rewrite each prompt using the six-part formula (subject · action · setting · camera · light & mood · audio).

Step 2 — Nano Banana for keyframes & ingredients

- **Pro for finals, NB2 for speed** — Nano Banana Pro renders up to 4K with reliable on-image typography; Nano Banana 2 is faster and grounded in live web data for accurate real-world detail.
- **Edit conversationally** — Don't restart on a near-miss. "Same image, but golden hour and remove the third person" preserves everything that already worked.
- **Build an ingredient kit** — Generate your character, location and key prop as separate images — these become the "ingredients" Veo combines into one coherent scene.

Step 3 — Veo 3.1 to animate

- **Ingredients-to-Video** — Feed up to three reference images plus a short prompt; Veo blends them with consistent identity across scenes — the single most useful feature for narrative work.
- **Keyframe bracketing** — Give Veo a first and last frame and let it generate the motion between them. Maximum control over where a shot starts and ends.
- **Go vertical natively** — Veo 3.1 outputs 9:16 directly and upscales to 1080p/4K — stop cropping horizontal masters for Reels and Shorts.
- **Write the dialogue in** — Veo generates speech and ambience natively. Put spoken lines in quotes inside the prompt and describe the soundscape explicitly.

Access points: gemini.google.com (app), Flow (Google's AI filmmaking studio), Google Vids, and the Gemini API/Vertex AI for automation. Workspace business accounts include Nano Banana Pro and Veo 3.1 on qualifying plans.

4 · Prompt Library — Steal These

Eight templates built on the six-part formula. Replace the bracketed parts; keep the structure.

The formula

[SUBJECT] + [ACTION] + [SETTING] + [CAMERA] + [LIGHT & MOOD] + [AUDIO]

News establishing shot

Aerial establishing shot of [city skyline] at dawn, slow forward drone push, soft golden light breaking through haze, documentary realism. Ambient city hum, distant traffic.

Interview b-roll

A [profession] works at [location], hands in close-up adjusting [object]. Shallow depth of field, slow lateral dolly, warm practical lighting, natural ambience with quiet room tone.

Product / object reveal

[Product] rotating on a dark reflective surface, macro lens, single hard key light sweeping across, slow 180-degree orbit, premium minimal mood. Subtle whoosh and low synth pad.

Character dialogue (Veo / Gen-4.5)

[Character description] looks into the lens and says: “[line of dialogue]”. Medium close-up, static tripod shot, soft window light, sincere tone. Clean dialogue, faint outdoor ambience.

Archive / period look

[Scene] shot on 16mm film, visible grain and gate weave, slightly desaturated, handheld, 1980s news-archive aesthetic. Optical soundtrack hiss, period-correct street sounds.

Social vertical hook

Vertical 9:16. [Subject] performs [surprising action] directly toward camera, fast push-in, high-contrast punchy color, energetic mood. Impact sound on the beat, trending-track-ready.

Image-to-video animation

Animate this frame: [describe ONLY the motion – e.g. “steam rises from the cup, her eyes blink, light flickers”]. Camera locked. Preserve composition, identity and lighting exactly.

Keyframe prompt (Nano Banana)

Cinematic still: [scene], 35mm lens, [light], color-graded for [mood]. Composition with clear negative space [left/right] for text. Photorealistic, 4K detail.

5 • Audio Generation Playbook

Half your production value is sound. The 2026 toolbox covers voice, music, SFX and transcription.

Voice & dialogue

- **ElevenLabs v3 is the broadcast standard** — Most expressive TTS on the market, ~70+ languages, with audio tags like [whispers], [laughs] or [sighs] to direct performance inside the script.
- **Multilingual v2 for stability** — When you need a long, consistent narration with zero surprises (documentaries, explainers), the older v2 model is still the safest choice.
- **Clone responsibly** — Voice cloning requires informed consent — non-negotiable in a newsroom. Document permission before you clone any real person's voice.

Native audio in video models

- **Let the model do the first pass** — Runway Gen-4.5 and Veo 3.1 generate dialogue, ambience and sync sound with the picture — often good enough for ambience, rarely final for dialogue.
- **Replace, don't fix** — If generated dialogue is 90% right, re-record that line with TTS and lay it over — faster than re-rolling the whole shot.

Music & SFX

- **Eleven Music / Lyria for score** — Text-to-music trained on licensed data — describe genre, tempo, mood and instrumentation, and request stems if you'll edit against picture.
- **Text-to-SFX for foley** — “Rain on a tin roof”, “newsroom keyboard clatter”, “crowd gasp” — describe it, get a file. Build a personal SFX library from your best generations.

Transcription & accessibility

- **Scribe v2 Realtime** — Low-latency live speech-to-text across dozens of languages — use it for live subtitling, interview logging and same-day multilingual cuts.
- **Always human-check Arabic output** — Dialect handling has improved dramatically, but broadcast Arabic still needs an editor's pass — treat AI transcription as a first draft.

6 • One-Page Cheat Sheet

Pin this above your edit bay.

DO	DON'T
Plan shots on paper before you spend a single credit	Treat generation like a slot machine — re-rolling without changing the prompt
Generate a keyframe image first, then animate it	Cram three scenes into one prompt
Use film language: dolly, rack focus, handheld, golden hour	Stack vague adjectives: "beautiful, amazing, cinematic, epic"
Save seeds and reference images from every winning shot	Publish AI footage of real events without labelling it
Disclose AI-generated content per your newsroom's policy	Clone a voice or likeness without written consent

Mini-glossary

Text-to-video (T2V) — Generating video purely from a written prompt.

Image-to-video (I2V) — Animating a still frame; the frame fixes composition and identity.

Reference / Ingredient — An image the model must stay consistent with across generations.

Keyframe — A first (and optionally last) frame the model animates between.

Seed — The random number behind a generation; same seed + same prompt ≈ same result.

Native audio — Sound generated together with the picture by the video model itself.

Motion brush — Painting a region of an image to control what moves and how.

Upscale — Raising output resolution after generation (e.g. to 1080p or 4K).

Stay in touch

Tag what you build: [@a.abudiak](#) on Instagram · connect on LinkedIn: [linkedin.com/in/abdullahabudiak](https://www.linkedin.com/in/abdullahabudiak).

Tomorrow — Day 2: deeper into the post-production pipeline. Bring today's best generation with you.